

Assignment 2 for Artificial Intelligence Question 1 Report

Name: Aishwarya Mahindru, RollNo.: g25ait1012

Github Link: <https://github.com/AishwaryaMahindru/nas-ga-basic>

1. Q1A – Roulette-Wheel Selection

In the original assignment template, parent selection was expected via tournament selection (commented out in your file). This was replaced with **Roulette-Wheel Selection**, implemented directly inside the `selection()` method.

- **Code Behaviour:**
 - Computes fitness values for all individuals:
$$\text{fitnesses} = [c.\text{fitness} \text{ for } c \text{ in self.population}]$$
 - Normalizes them to obtain probabilities.
 - Builds a cumulative probability list.
 - Draws two random values to select two indices:
$$\text{selected.append(idx)}$$
 - These indices are later converted to actual parents inside `evolve()`:
$$\text{parent1} = \text{self.population}[idx1]$$

$$\text{parent2} = \text{self.population}[idx2]$$
- **How this is valid:**
 - Higher-fitness architectures get proportionally higher chances.
 - Lower-fitness ones still retain some probability → preserves genetic diversity.
 - Matches biologically-inspired evolutionary selection better than tournament selection.
- **Code changes:**
 - Removed tournament selection (commented out).
 - Added roulette-wheel logic inside `selection()`.
 - Updated `evolve()` to correctly convert selected indices into
$$\text{parent1} = \text{self.population}[idx1]$$

$$\text{parent2} = \text{self.population}[idx2]$$
 - Ensured crossover and mutation receive full Architecture objects.
- **NAS Run Log Github link:** https://github.com/AishwaryaMahindru/nas-ga-basic/blob/main/outputs/run_1/nas_run.log

2. Q1B – Modified Fitness Function (Weighted Conv/FC Penalty)

I revised the fitness computation inside `evaluate_fitness()` by introducing a weighted parameter penalty based on architecture complexity.

- **What the Code Computes**
 - After training and obtaining `best_acc`, it computes two parameter counts:

- conv_params = 0
 - fc_params = 0
- ```
for name, param in model.named_parameters():
 if 'feature_extractor' in name:
 conv_params += param.numel()
 elif 'classifier' in name or 'head' in name:
 fc_params += param.numel()
```
- Then apply weights:  
 $\text{Conv weight: } 0.7$   
 $\text{FC weight: } 0.3$
  - Penalty:  
 $\text{penalty} = (w_{\text{conv}} * \text{conv\_params} + w_{\text{fc}} * \text{fc\_params}) * 1e-6$
  - Final fitness:  
 $\text{architecture.fitness} = \text{best\_acc} - \text{penalty}$
- **Justification of Weights:**
    - Convolution parameters dominate computational cost (most FLOPs).
    - Conv weights are reused across spatial locations → very expensive.
    - FC layers contribute fewer multiply-accumulate ops → lower cost.
    - Penalizing conv parameters more avoids evolution of overly large conv blocks.
    - Lower FC penalty preserves classifier capacity.
    - (0.7, 0.3) resembles real CNN compute distribution.

### 3. Results:

#### Generation 1 Results

| Arch | Conv Params | FC Params | Old Penalty | New Penalty | Diff     | Fitness        | Acc  |
|------|-------------|-----------|-------------|-------------|----------|----------------|------|
| 1    | 1,216       | 10,51,402 | 1.05265     | 0.31627     | -0.73638 | 0.23873        | 0.56 |
| 2    | 27,536      | 21,02,794 | 2.13043     | 0.65011     | -1.48031 | -0.04411       | 0.61 |
| 3    | 2,41,312    | 10,49,994 | 1.29188     | 0.48392     | -0.80797 | 0.15208        | 0.64 |
| 4    | 1,792       | 41,97,130 | 4.19905     | 1.26039     | -2.93866 | -0.73039       | 0.53 |
| 5    | 19,888      | 1,31,786  | 0.15183     | 0.05346     | -0.09838 | <b>0.55254</b> | 0.61 |
| 6    | 2,03,264    | 41,95,722 | 4.39931     | 1.401       | -2.9983  | -0.882         | 0.52 |
| 7    | 28,032      | 10,49,290 | 1.07748     | 0.33441     | -0.74307 | 0.18459        | 0.52 |
| 8    | 15,248      | 21,02,794 | 2.11814     | 0.64151     | -1.47663 | -0.05251       | 0.59 |
| 9    | 2,432       | 41,99,946 | 4.20244     | 1.26169     | -2.94076 | -0.73069       | 0.53 |
| 10   | 2,432       | 20,99,978 | 2.10247     | 0.6317      | -1.47078 | -0.0527        | 0.58 |

#### Generation 2 Results

| Arch | Conv Params | FC Params | Old Penalty | New Penalty | Diff     | Fitness        | Acc  |
|------|-------------|-----------|-------------|-------------|----------|----------------|------|
| 1    | 19,888      | 1,31,786  | 0.15183     | 0.05346     | -0.09838 | <b>0.52454</b> | 0.58 |
| 2    | 1,216       | 10,51,402 | 1.05265     | 0.31627     | -0.73638 | 0.21373        | 0.53 |
| 3    | 2,432       | 41,99,946 | 4.20244     | 1.26169     | -2.94076 | -0.70469       | 0.56 |

|    |       |           |         |         |          |          |      |
|----|-------|-----------|---------|---------|----------|----------|------|
| 4  | 1,792 | 41,97,130 | 4.19905 | 1.26039 | -2.93866 | -0.73339 | 0.53 |
| 5  | 448   | 10,51,402 | 1.05188 | 0.31573 | -0.73615 | 0.22627  | 0.54 |
| 6  | 2,432 | 41,99,946 | 4.20244 | 1.26169 | -2.94076 | -0.72469 | 0.54 |
| 7  | 1,792 | 41,97,130 | 4.19905 | 1.26039 | -2.93866 | -0.74139 | 0.52 |
| 8  | 1,792 | 41,97,130 | 4.19905 | 1.26039 | -2.93866 | -0.74339 | 0.52 |
| 9  | 1,216 | 21,02,794 | 2.10404 | 0.63169 | -1.47235 | -0.09969 | 0.53 |
| 10 | 2,432 | 41,99,946 | 4.20244 | 1.26169 | -2.94076 | -0.72169 | 0.54 |

### Generation 3 Results

| Arch | Conv Params | FC Params | Old Penalty | New Penalty | Diff     | Fitness        | Acc  |
|------|-------------|-----------|-------------|-------------|----------|----------------|------|
| 1    | 19,888      | 1,31,786  | 0.15183     | 0.05346     | -0.09838 | <b>0.55854</b> | 0.61 |
| 2    | 448         | 10,51,402 | 1.05188     | 0.31573     | -0.73615 | 0.20627        | 0.52 |
| 3    | 1,216       | 21,02,794 | 2.10404     | 0.63169     | -1.47235 | -0.10269       | 0.53 |
| 4    | 1,216       | 21,02,794 | 2.10404     | 0.63169     | -1.47235 | -0.11269       | 0.52 |
| 5    | 1,216       | 21,02,794 | 2.10404     | 0.63169     | -1.47235 | -0.08369       | 0.55 |
| 6    | 1,792       | 41,97,130 | 4.19905     | 1.26039     | -2.93866 | -0.70239       | 0.56 |
| 7    | 45,968      | 20,99,978 | 2.14636     | 0.66217     | -1.48419 | -0.04517       | 0.62 |
| 8    | 4,05,520    | 2,64,970  | 0.67116     | 0.36336     | -0.30781 | 0.30765        | 0.67 |
| 9    | 1,792       | 41,97,130 | 4.19905     | 1.26039     | -2.93866 | -0.72739       | 0.53 |
| 10   | 9,728       | 83,91,434 | 8.40142     | 2.52424     | -5.87718 | -2.03624       | 0.49 |

### Generation 4 Results

| Arch | Conv Params | FC Params | Old Penalty | New Penalty | Diff     | Fitness        | Acc  |
|------|-------------|-----------|-------------|-------------|----------|----------------|------|
| 1    | 19,888      | 1,31,786  | 0.15183     | 0.05346     | -0.09838 | <b>0.55854</b> | 0.61 |
| 2    | 4,05,520    | 2,64,970  | 0.67116     | 0.36336     | -0.30781 | 0.30465        | 0.67 |
| 3    | 2,432       | 20,99,978 | 2.10247     | 0.6317      | -1.47078 | -0.0487        | 0.58 |
| 4    | 1,216       | 21,02,794 | 2.10404     | 0.63169     | -1.47235 | -0.10269       | 0.53 |
| 5    | 1,792       | 41,97,130 | 4.19905     | 1.26039     | -2.93866 | -0.71039       | 0.55 |
| 6    | 1,792       | 41,97,130 | 4.19905     | 1.26039     | -2.93866 | -0.71839       | 0.54 |
| 7    | 1,792       | 41,97,130 | 4.19905     | 1.26039     | -2.93866 | -0.69539       | 0.57 |
| 8    | 1,792       | 41,97,130 | 4.19905     | 1.26039     | -2.93866 | -0.72239       | 0.54 |
| 9    | 69,920      | 10,54,218 | 1.12436     | 0.36521     | -0.75915 | 0.25479        | 0.62 |
| 10   | 1,216       | 21,02,794 | 2.10404     | 0.63169     | -1.47235 | -0.09969       | 0.53 |

### Generation 5 Results

| Arch | Conv Params | FC Params | Old Penalty | New Penalty | Diff     | Fitness        | Acc  |
|------|-------------|-----------|-------------|-------------|----------|----------------|------|
| 1    | 19,888      | 1,31,786  | 0.15183     | 0.05346     | -0.09838 | <b>0.56254</b> | 0.62 |
| 2    | 4,05,520    | 2,64,970  | 0.67116     | 0.36336     | -0.30781 | 0.31865        | 0.68 |
| 3    | 4,05,520    | 2,64,970  | 0.67116     | 0.36336     | -0.30781 | 0.31465        | 0.68 |
| 4    | 4,05,520    | 2,64,970  | 0.67116     | 0.36336     | -0.30781 | 0.29265        | 0.66 |
| 5    | 4,05,520    | 2,64,970  | 0.67116     | 0.36336     | -0.30781 | 0.30865        | 0.67 |
| 6    | 1,216       | 21,02,794 | 2.10404     | 0.63169     | -1.47235 | -0.08469       | 0.55 |
| 7    | 4,05,520    | 2,64,970  | 0.67116     | 0.36336     | -0.30781 | 0.32065        | 0.68 |

|    |          |           |         |         |          |          |      |
|----|----------|-----------|---------|---------|----------|----------|------|
| 8  | 4,05,520 | 2,64,970  | 0.67116 | 0.36336 | -0.30781 | 0.30765  | 0.67 |
| 9  | 1,216    | 21,02,794 | 2.10404 | 0.63169 | -1.47235 | -0.09569 | 0.54 |
| 10 | 1,216    | 21,02,794 | 2.10404 | 0.63169 | -1.47235 | -0.08169 | 0.55 |

#### Summary of Best Architecture per Generation

| Generation | Conv Layers | Best Fitness   | Accuracy     |
|------------|-------------|----------------|--------------|
| Gen-1      | 3           | 0.55254        | 0.606        |
| Gen-2      | 3           | 0.52454        | 0.578        |
| Gen-3      | 3           | 0.55854        | 0.612        |
| Gen-4      | 3           | 0.55854        | 0.612        |
| Gen-5      | 3           | <b>0.56254</b> | <b>0.616</b> |

#### Final Best Architecture

| Feature        | Value            |
|----------------|------------------|
| Conv layers    | <b>3</b>         |
| Conv configs   | 32-5, 16-3, 32-5 |
| Pooling        | Max              |
| Activation     | ReLU             |
| FC Units       | 64               |
| Final Accuracy | <b>0.616</b>     |
| Final Fitness  | <b>0.5625</b>    |
| Total Params   | 1,51,834         |

#### Roulette Wheel Selection Details:

| Index    | Fitness      | Probability | Cumulative |
|----------|--------------|-------------|------------|
| <b>0</b> | 0.56254      | 0.260034    | 0.260034   |
| <b>1</b> | 0.32065      | 0.148217    | 0.408251   |
| <b>2</b> | 0.31865      | 0.147293    | 0.555544   |
| <b>3</b> | 0.31465      | 0.145444    | 0.700988   |
| <b>4</b> | 0.30865      | 0.14267     | 0.843658   |
| <b>5</b> | 0.30765      | 0.142208    | 0.985866   |
| <b>6</b> | 0.29265      | 0.135274    | 1.12114    |
| <b>7</b> | -<br>0.08169 | -0.037761   | 1.08338    |
| <b>8</b> | -<br>0.08469 | -0.039147   | 1.044232   |
| <b>9</b> | -<br>0.09569 | -0.044232   | 1          |

**High-Level Evolution Summary Table**

| Generation | Dominant Conv Depth             | Common Filters | Common Activation | Pool    | Notes                  |
|------------|---------------------------------|----------------|-------------------|---------|------------------------|
| 0          | 1 conv                          | 16, 32, 64     | ReLU + Leaky      | avg     | Very simple networks   |
| 1          | 1 + some 3/4 conv               | 16, 32, 64     | ReLU              | avg/max | More variety           |
| 2          | Mixed 1, 3, 4 conv              | 32, 64, 128    | ReLU              | max+avg | Complex filters emerge |
| 3          | <b>4 conv dominant</b>          | 64, 128        | ReLU              | avg     | Convergence begins     |
| 4          | <b>4 conv strongly dominant</b> | 64, 128        | ReLU              | avg     | Population converged   |

#### Interpretation:

- The GA started with mostly 1-conv shallow models in Generation 0.
- As evolution progressed, deeper 3-conv and 4-conv architectures appeared and gained higher fitness.
- By Generations 3–4, the population strongly converged to 4-conv CNNs, indicating they performed best.
- ReLU became the dominant activation since it consistently gave higher accuracy.
- Avg pooling emerged as the stable pooling choice in later generations.
- FC units converged to 256, showing the GA found this configuration most efficient and reliable.

#### 4. Learnings & Reflections

- Understanding Genetic Algorithms in NAS
  - I learned how Genetic Algorithms can guide neural architecture search by evolving CNN structures generation after generation.
  - I saw firsthand how even a small change in selection, mutation, or fitness calculation affects the entire search trajectory.
- Debugging & Mistakes Identified
  - I initially assumed convolution layers would contain "conv" in their parameter names, which caused conv parameters to always be zero.
    - This was fixed by switching to module-type checks (`nn.Conv2d`, `nn.Linear`) instead of name-based checks.
  - Another mistake was selecting parent indices and directly passing them into crossover/mutation.
    - This caused attribute errors until I correctly mapped indices back to architecture objects.
  - Throughout the assignment, I made multiple Git commits fixing mistakes step-by-step, which helped me track changes clearly and understand the evolution of my implementation.
- Importance of Logging & Iterative Reruns
  - Adding logging for fitness probabilities, parameter counts, and penalty differences made debugging much easier.
  - Multiple reruns taught me that evolutionary algorithms are extremely sensitive — even

small bugs drastically change the best architecture found.

➤ Training Efficiency & Hardware Impact

- I realized how costly CNN evaluation inside a GA is, since every architecture must be trained, even if briefly.
- Using an **NVIDIA A100 GPU** made a huge difference:
  - Runtime that took hours on CPU completed in minutes.
  - Faster feedback loops helped me debug and iterate much more effectively.

## 5. System Used and Runtime

- GPU – A100 (via Google Colab)
- Runtime – 47 minutes
- Snapshot of GPU Resource Utilization: [https://github.com/AishwaryaMahindru/nas-ga-basic/blob/main/outputs/GPU\\_Run\\_Resources.png](https://github.com/AishwaryaMahindru/nas-ga-basic/blob/main/outputs/GPU_Run_Resources.png)

```
trainset = CIFAR10(root='./data', train=True, download=True, transform=transform)
valset = CIFAR10(root='./data', train=False, download=True, transform=transform)

Use only 5000 samples for quick NAS
train_subset = Subset(trainset, range(5000))
val_subset = Subset(valset, range(1000))

train_loader = DataLoader(train_subset, batch_size=256, shuffle=True)
val_loader = DataLoader(val_subset, batch_size=256, shuffle=False)

Run NAS with GA
ga = GeneticAlgorithm(
 population_size=10, # Small population for demonstration
 generations=5, # Few generations for quick results
 mutation_rate=0.3,
 crossover_rate=0.7
)

best_arch = ga.evolve(train_loader, val_loader, device, run=len(all_log))

print(f"\n{'='*60}", flush=True)
print("FINAL BEST ARCHITECTURE", flush=True)
print(f"{'='*60}", flush=True)
print(f"Genes: {best_arch.genes}", flush=True)
print(f"Accuracy: {best_arch.accuracy:.4f}", flush=True)
print(f"Fitness: {best_arch.fitness:.4f}", flush=True)

Build and test final model
final_model = CNN(best_arch.genes).to(device)
```